

SHUBHAM KUMAR

B.Tech Computer Science · Data Science & AI

Delhi, India | [+91 73618 57543](tel:917361857543) | yshubham0103@gmail.com | linkedin.com/in/shubham-kumar-288b7437b | github.com/shubham00011222 | [Portfolio](#)

SUMMARY

B.Tech CS student at NIT Delhi (2023–27) with hands-on experience building end-to-end ML pipelines, SQL analytics systems, and NLP applications across telecom, finance, and sports domains. Proficient in Python, SQL, XGBoost, LightGBM, and Streamlit. All projects use real public datasets with honest, verifiable metrics. Seeking a Data Science or Data Analytics internship.

EDUCATION

National Institute of Technology Delhi (NIT Delhi)

Delhi, India | 2023 – 2027

B.Tech in Computer Science & Engineering

TECHNICAL SKILLS

Languages: Python, SQL

ML / DL: Scikit-learn, XGBoost, LightGBM, PyTorch, Keras, SHAP, EfficientNet

NLP & LLMs: HuggingFace Transformers, FinBERT, spaCy, NLTK

Data Stack: Pandas, NumPy, SQLite

Visualisation: Streamlit, Plotly, Seaborn, Matplotlib

MLOps & Deployment: Docker, FastAPI, PostgreSQL, Streamlit Cloud, HuggingFace Spaces, Vercel

PROJECTS

IPL Analytics (2008–2024) — End-to-End Cricket Analytics

[GitHub](#) [Live Demo](#)

Python · SQL (SQLite) · Pandas · Scikit-learn · Matplotlib · Seaborn · Streamlit

- Analysed 300K+ ball-by-ball IPL deliveries using 20+ SQL window function queries (CTEs, ROW_NUMBER, LAG) across batting, bowling, and team strategy. Built KMeans clustering to classify batsmen as Anchors vs Aggressors; logistic regression win-probability model (64.7% accuracy); 20 publication-quality visualisations including SHAP-style matchup heatmaps and phase-wise strike rate analysis. 8-page Streamlit dashboard with match simulator and rivalry deep-dive.

Telco Customer Churn Prediction Engine

[GitHub](#) [Live Demo](#)

Python · XGBoost · LightGBM · SHAP · FastAPI · Streamlit · Docker

- Trained XGBoost + LightGBM ensemble on IBM Telco Customer Churn dataset (7,043 real records); achieved ROC-AUC 0.717 with threshold-tuned recall of 64.2% on held-out test set. Applied SMOTE for class imbalance, SHAP TreeExplainer for per-customer reason codes, served via FastAPI (<50ms).

NLP Financial Sentiment Analyser

[GitHub](#) [Live Demo](#)

Python · FinBERT · HuggingFace Transformers · Streamlit

- Fine-tuned FinBERT (HuggingFace) on Financial PhraseBank dataset; achieved 87.3% accuracy and 0.86 macro F1 across bullish/bearish/neutral classes. Built automated ETL pipeline (NewsAPI → sentiment scoring) and live Streamlit dashboard tracking sentiment trends across 500+ tickers.

Retail Demand Forecasting System

[GitHub](#) [Live Demo](#)

Python · Prophet · LSTM · Keras · Streamlit · Docker

- Built Prophet + LSTM hybrid on Kaggle Store Sales dataset (2.6M rows, 54 stores, 33 product families); reduced MAPE from 22% to 9% vs naive baseline using holiday and promotion regressors. Generated 30-day forecasts with 90% confidence intervals across 1,200 SKUs.

Medical Image Classifier — Chest X-Ray Diagnosis

[GitHub](#) [Live Demo](#)

Python · PyTorch · EfficientNet-B4 · Grad-CAM · FastAPI · Streamlit

- Fine-tuned EfficientNet-B4 on NIH ChestX-ray14 (112K images); achieved 91.2% accuracy across 14 pathology classes. Applied Grad-CAM for visual explainability; deployed via FastAPI + Streamlit.

RELEVANT COURSEWORK

Data Structures & Algorithms · Database Management Systems · Probability & Statistics · Linear Algebra · Operating Systems · Computer Networks